



# Zero downtime on Kubernetes with Hazelcast

Nicolas Fränkel

# Me, myself and I

- Developer
- Developer Advocate
- We live in interesting times and I'm curious



@nicolas\_frankel

# Hazelcast



**HAZELCAST IMDG** is an **operational, in-memory**, distributed computing platform that manages data using in-memory storage and performs parallel execution for breakthrough application speed and scale.



**HAZELCAST JET** is the ultra fast, application embeddable, 3<sup>rd</sup> generation stream processing engine for low latency batch and stream processing.



@nicolas\_frankel

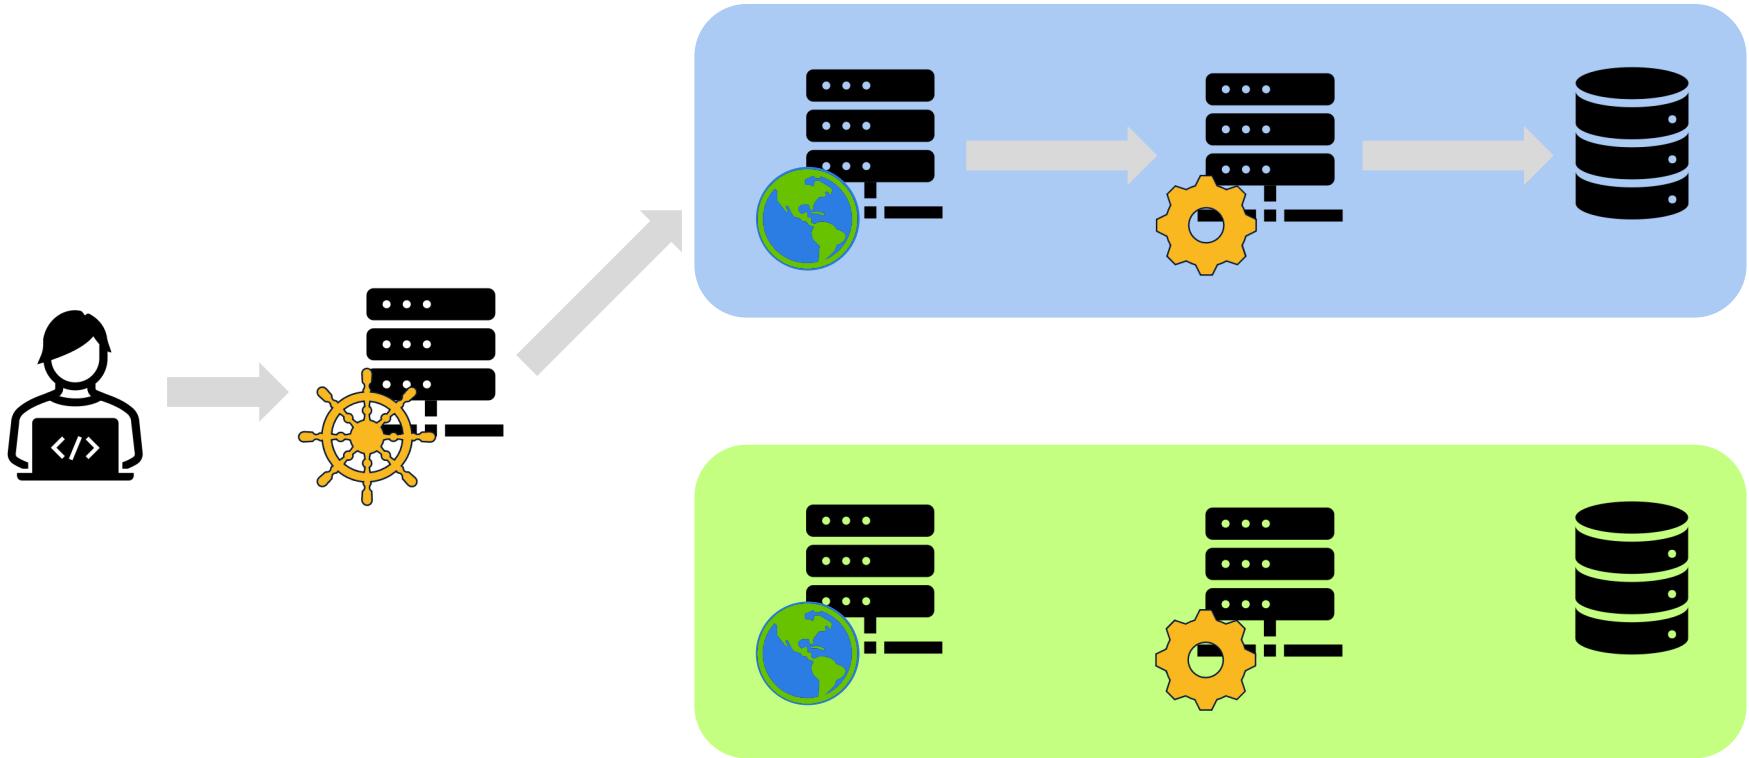
# Why zero downtime?

1. Business wants it
  - Downtime has a cost
2. Users expect it
  - When was the last time you saw Google Search display “Please come back later”?



@nicolas\_frankel

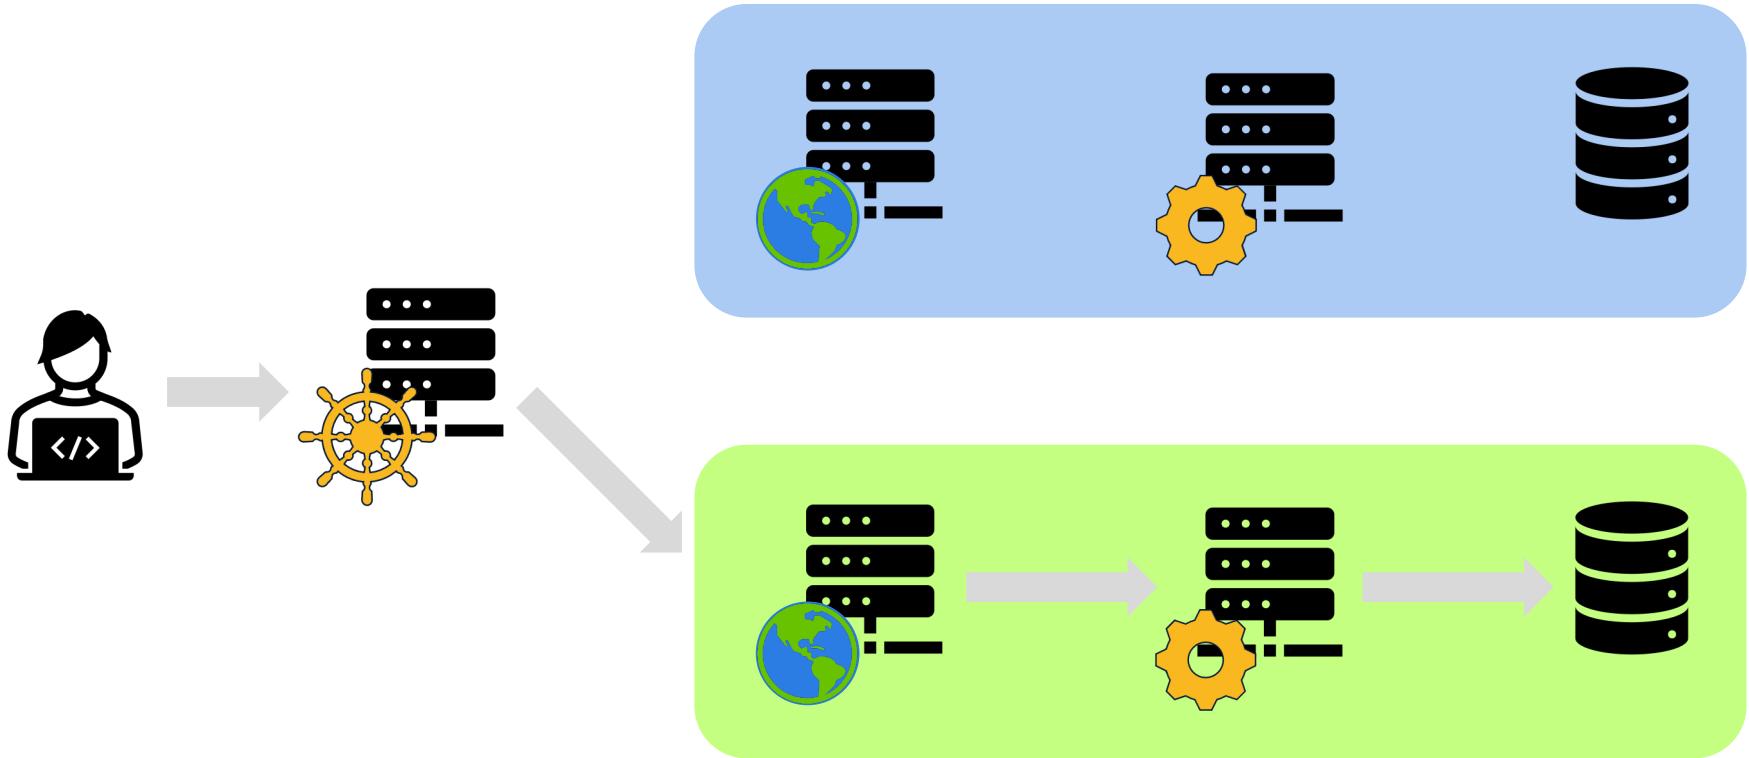
# Blue-Green deployment



@nicolas\_frankel



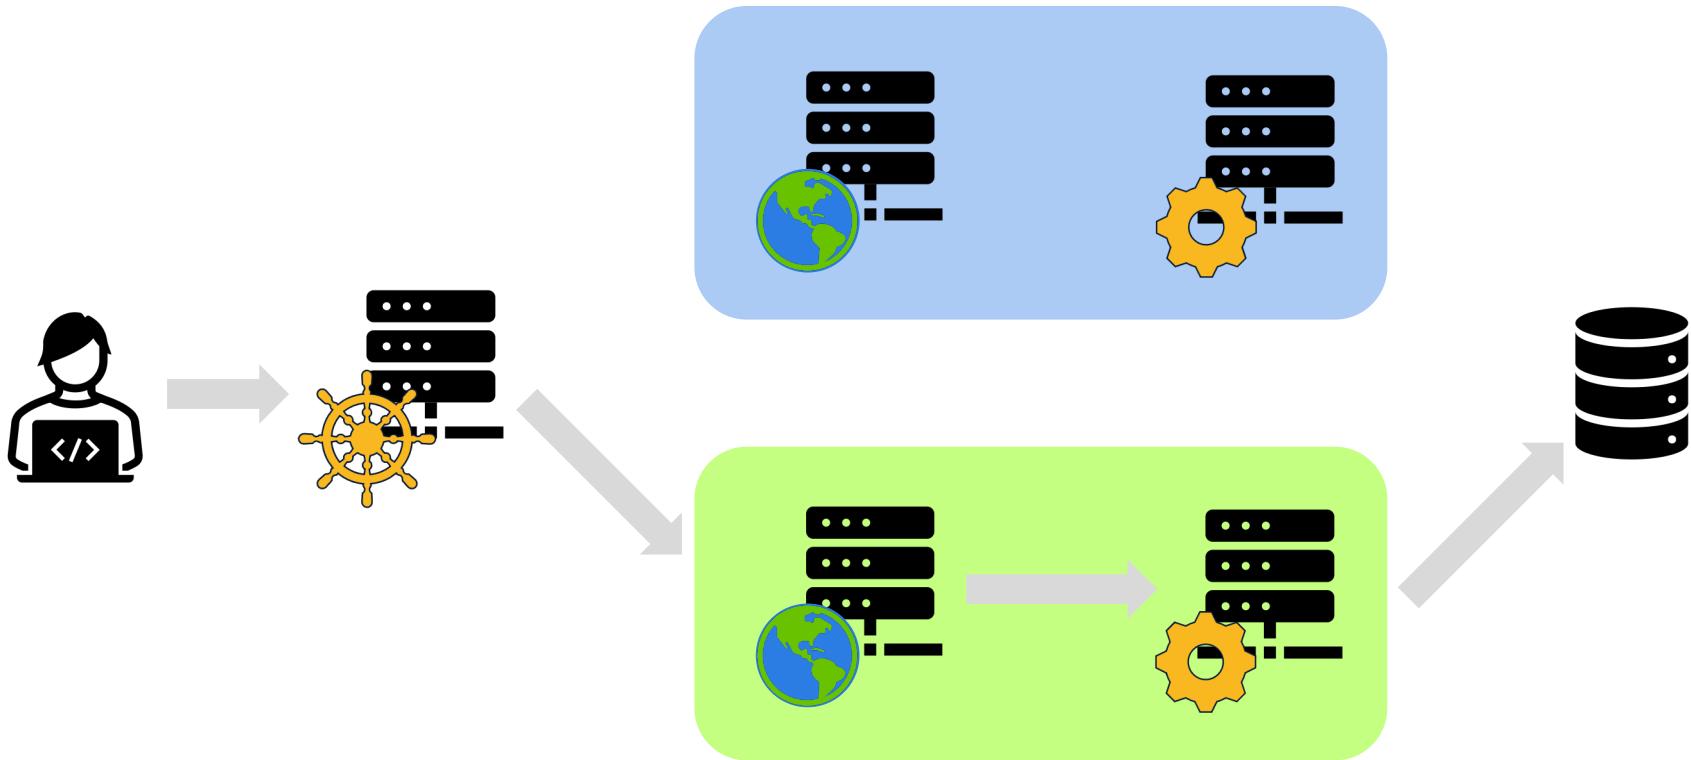
# Blue-Green deployment



@nicolas\_frankel



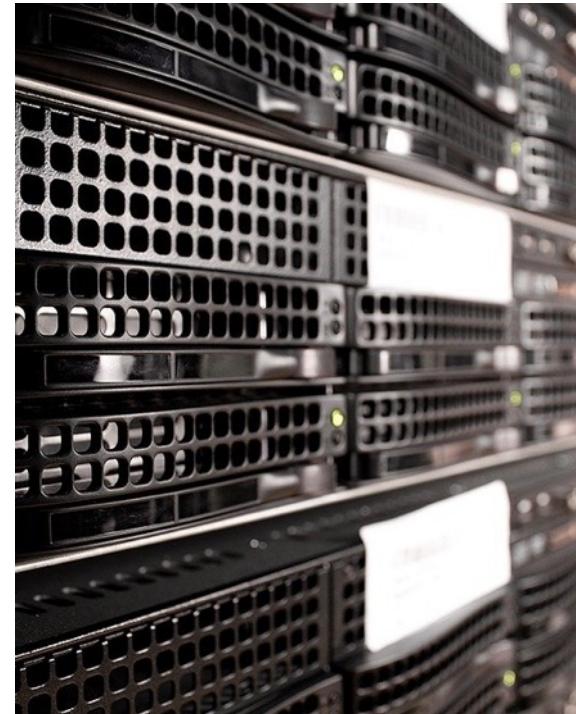
# Blue-Green deployment variant



@nicolas\_frankel

# Zero downtime's issues relate to state

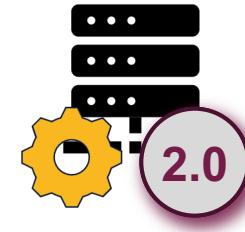
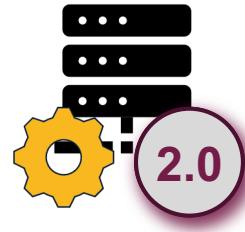
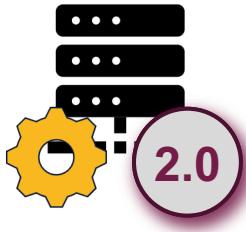
- State **in memory**
  - User sessions
- State **in the database**



@nicolas\_frankel



# Kubernetes rolling updates principle



@nicolas\_frankel



# Solving state-related issues

- Sessions
  - Session replication
- Database
  - That's the hard spot!



@nicolas\_frankel



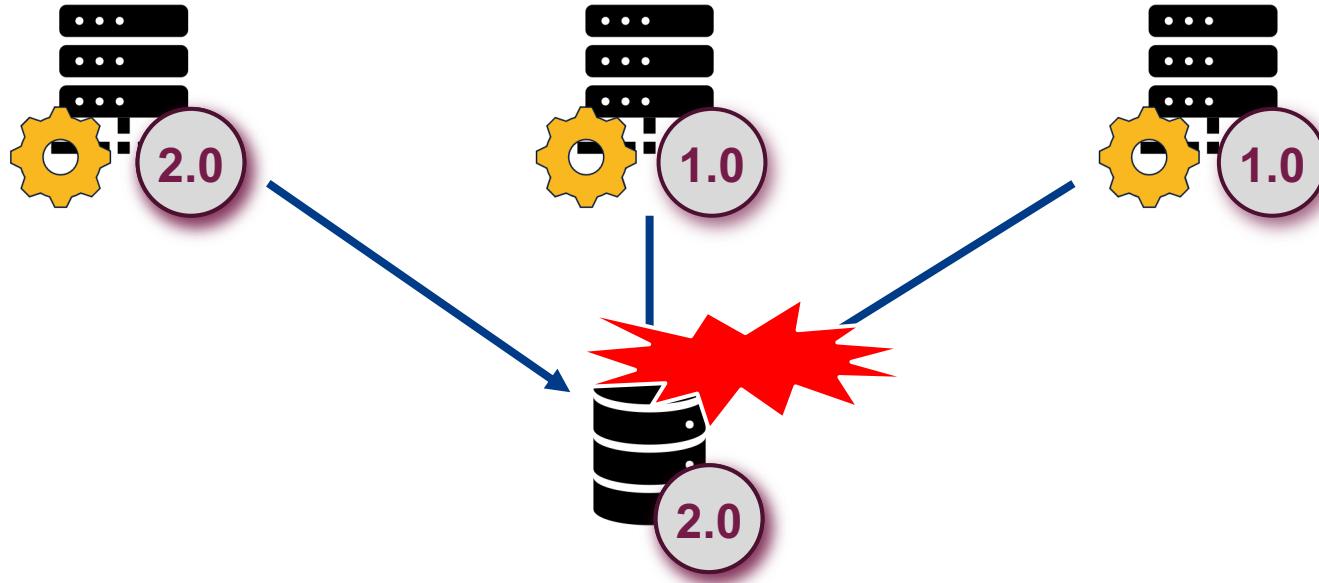
# Option 1: Keep the same database

- The application needs to cope with two versions of the schema



@nicolas\_frankel

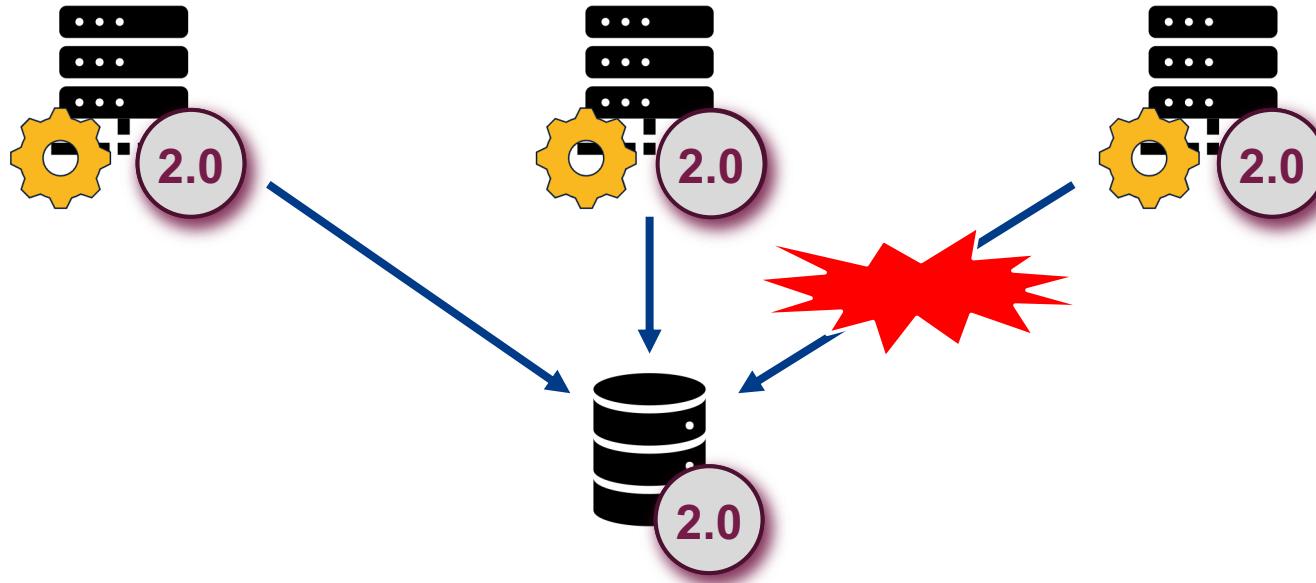
# Rolling upgrade issue with a single database



@nicolas\_frankel



# More issues with rollback



@nicolas\_frankel



# An e-commerce use-case



Powered by yFiles

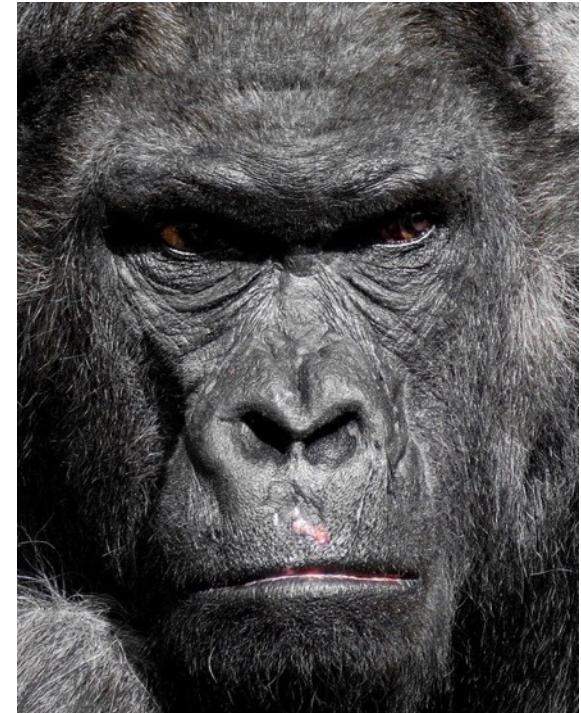


@nicolas\_frankel



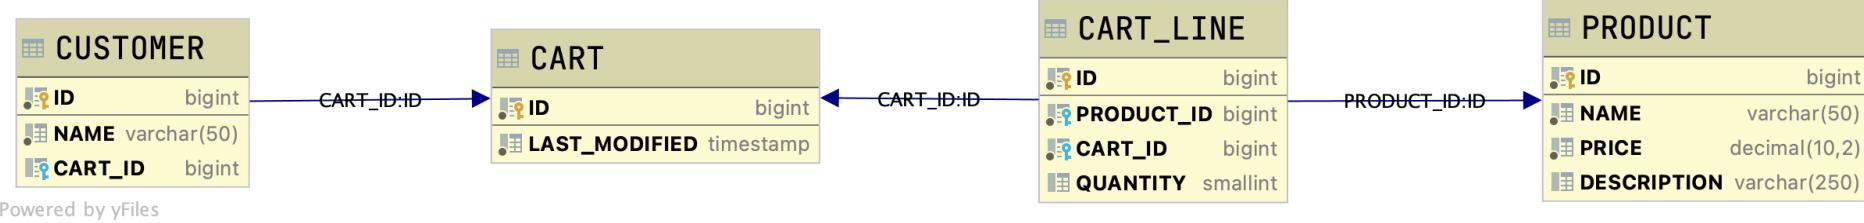
# New business requirement comes in!

- We want to keep a track when a cart was created
  - To send a reminder email after some time has passed



@nicolas\_frankel

# Target schema



@nicolas\_frankel



# Handling schema-related breaking changes

- Split the breaking change into a series of changes compatible side-by-side
- Plan for rollback (it happens!)



@nicolas\_frankel

# Steps' decomposition

1. Create CART table
  - App uses “old” data model
  - Trigger inserts CART when the first CART\_LINE is inserted
2. CART becomes the “source of truth”
  - App uses the CART table
  - Trigger updates CART\_LINE with CUSTOMER\_ID every time it’s inserted
3. Migration of untouched data
4. Cleanup

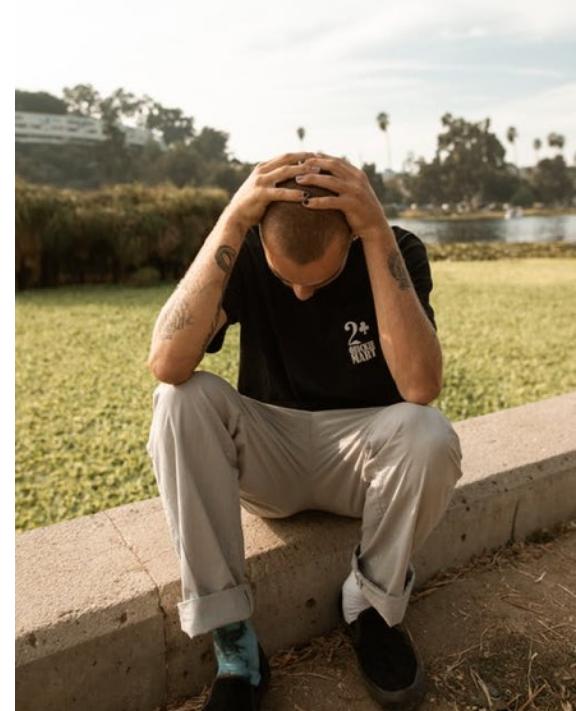


@nicolas\_frankel



# Issues of keeping the same database

- Requires steps' decomposition
- Rollback a single step only
- Needs planning across the organization (devs, DBAs, Ops)
- You will **need to migrate data anyway**



@nicolas\_frankel



# Option 2: Embrace data migration

- Have two different databases
- Migration implemented by:
  - Change-Data-Capture
  - **Data streaming**
- Developers are not impacted by Ops' concerns
- It works with any deployment option e.g. canary release



@nicolas\_frankel

# Change-Data-Capture

“In databases, Change Data Capture is a set of software design patterns used to **determine and track the data that has changed** so that action can be taken using the changed data.

CDC is an approach to data integration that is based on the **identification, capture and delivery of the changes made to enterprise data sources**.”

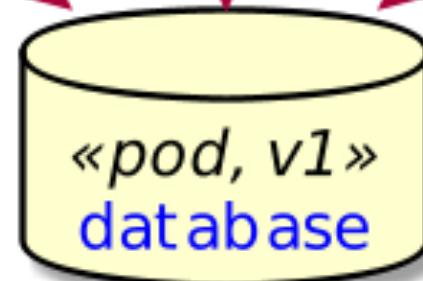
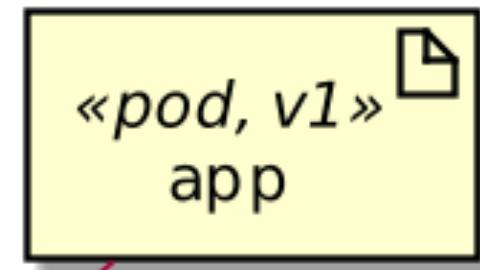
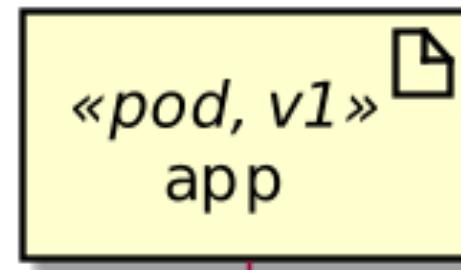
-- [https://en.wikipedia.org/wiki/Change\\_data\\_capture](https://en.wikipedia.org/wiki/Change_data_capture)



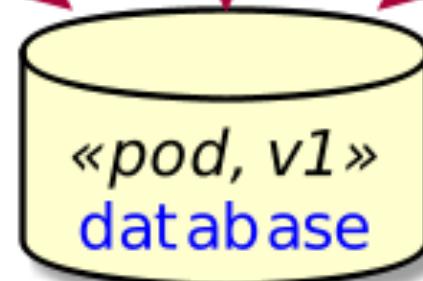
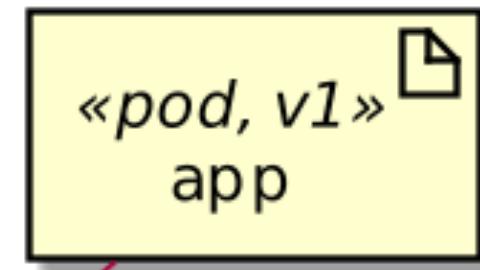
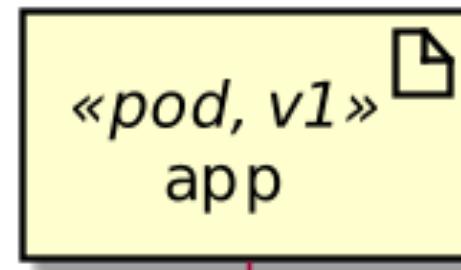
@nicolas\_frankel



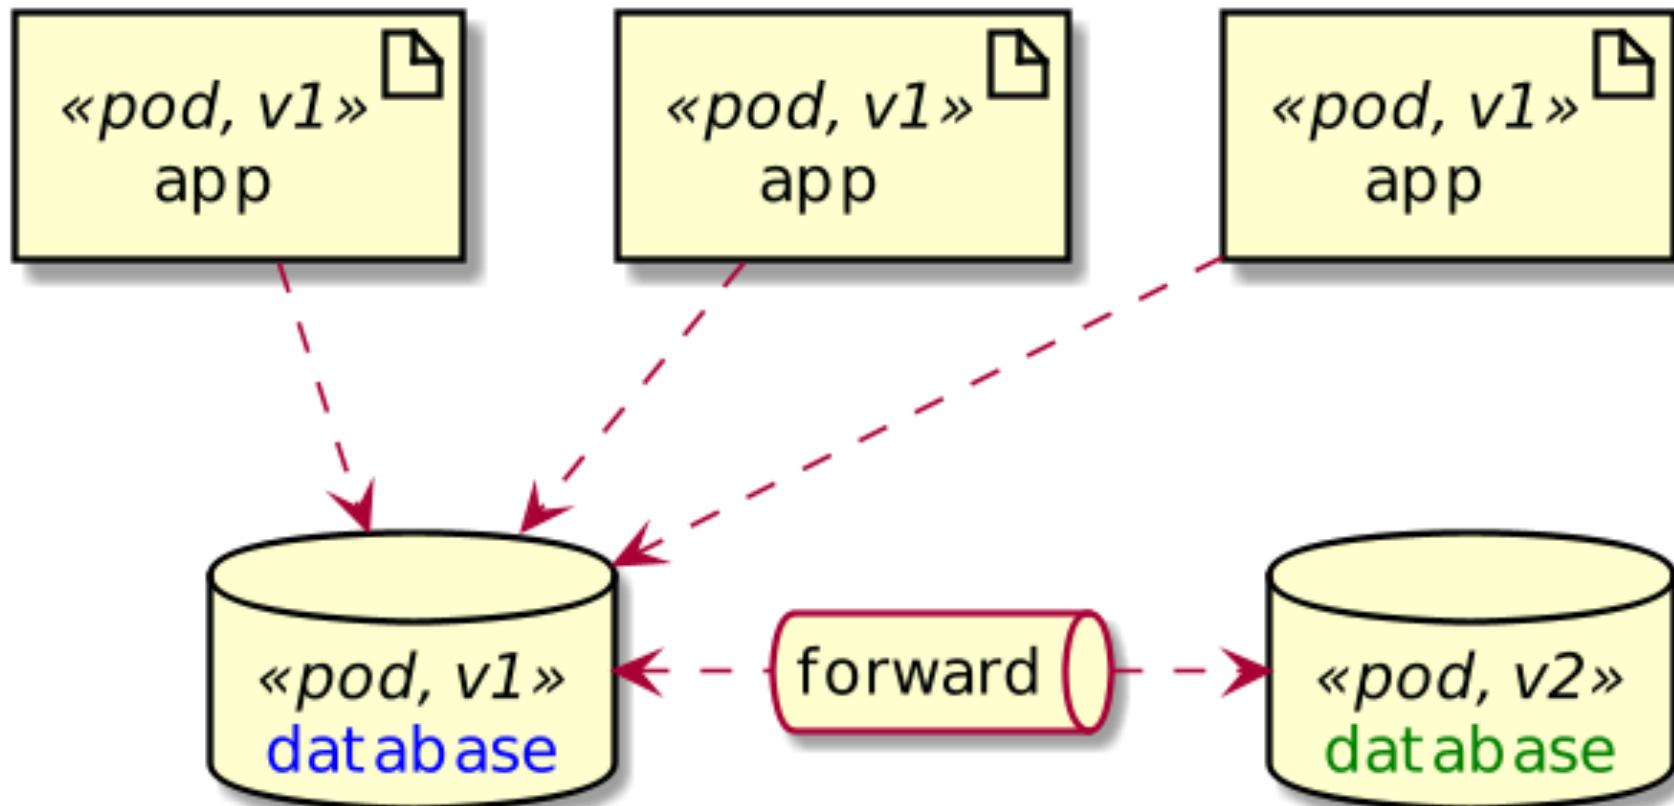
# Node



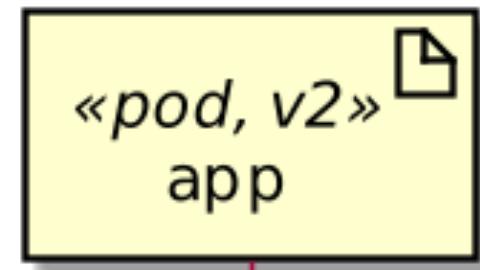
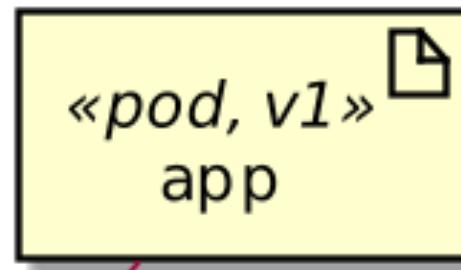
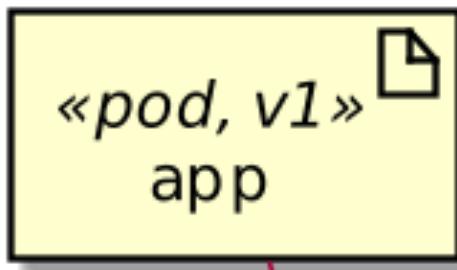
# Node



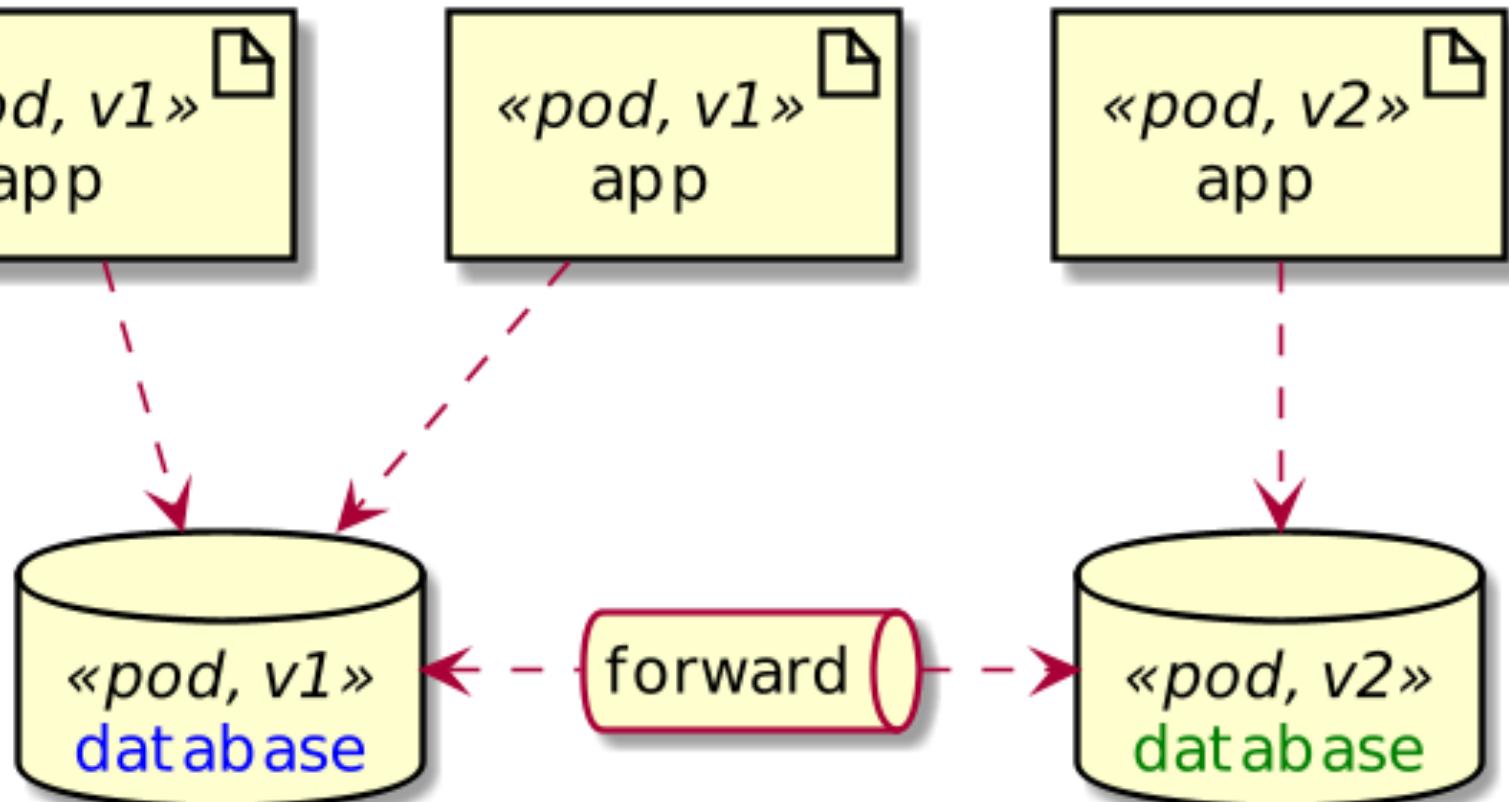
# Node



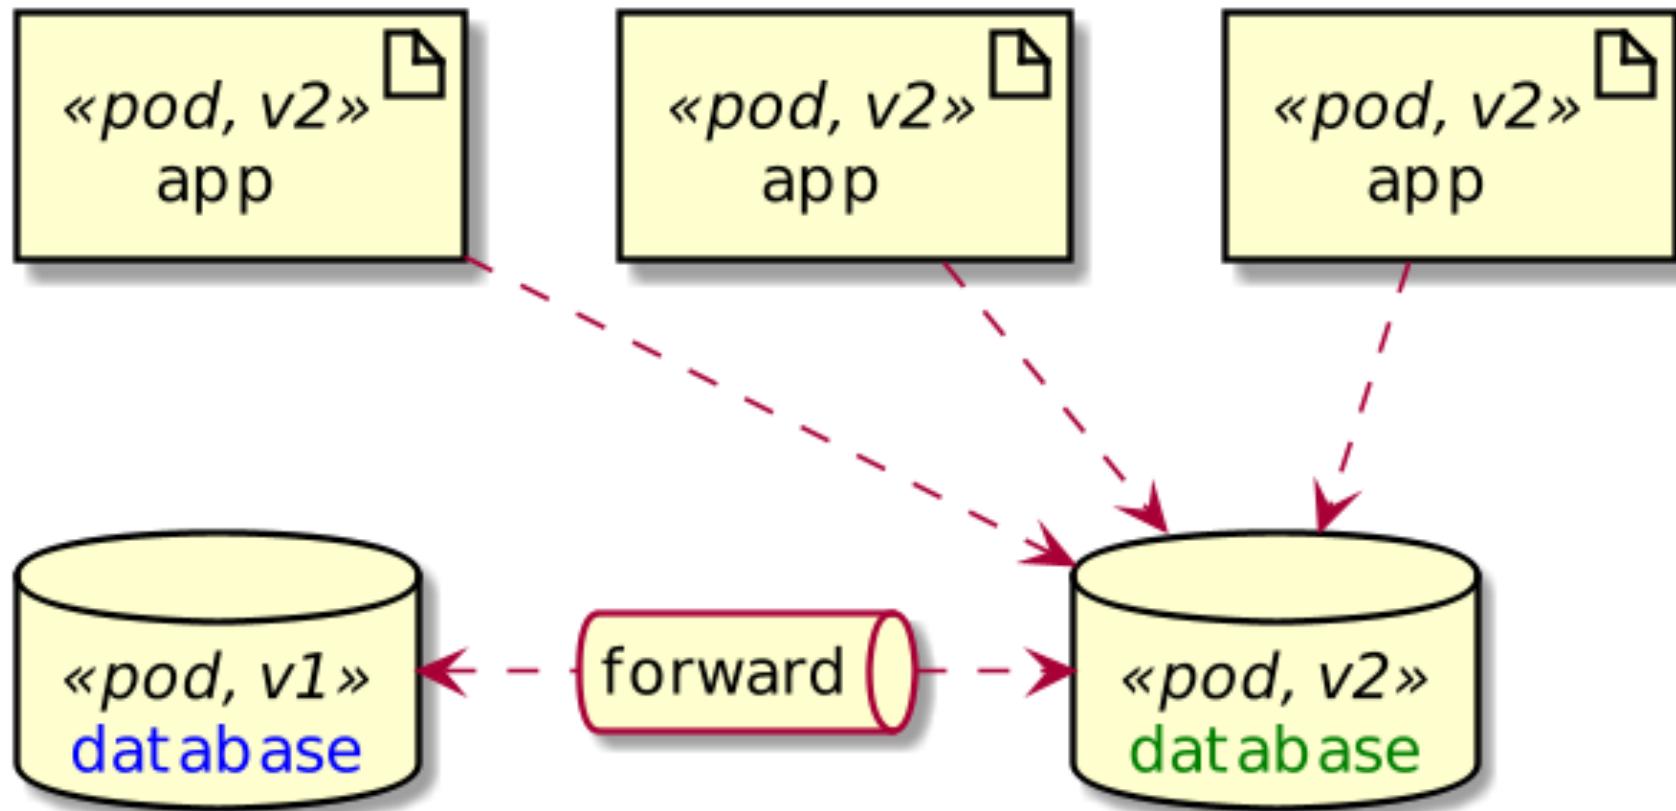
# Node



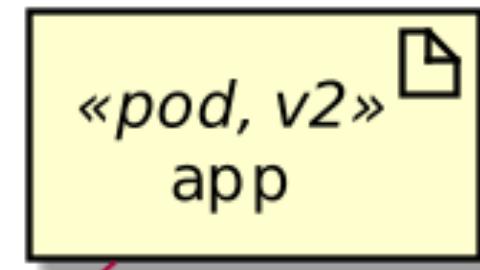
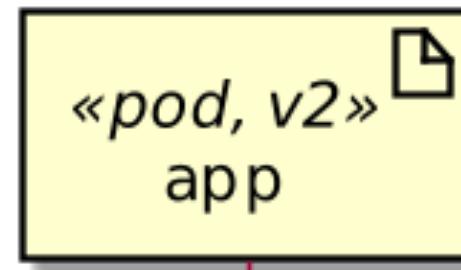
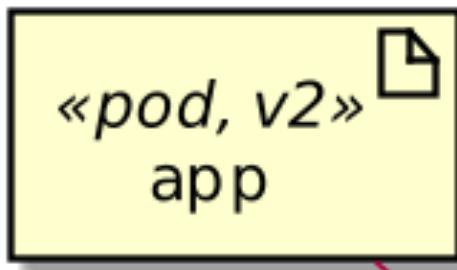
forward



# Node



# Node



# Implementation details

- Hazelcast for Session Replication
  - Via Spring Session
- Hazelcast for CDC
  - With Debezium



**hazelcast**<sup>®</sup>



@nicolas\_frankel



# Hazelcast IMDG

*“Use Hazelcast IMDG to store your data in RAM, spread and replicate it across a cluster of machines, and perform data-local computation on it. **Replication gives you resilience to failures of cluster nodes.**”*



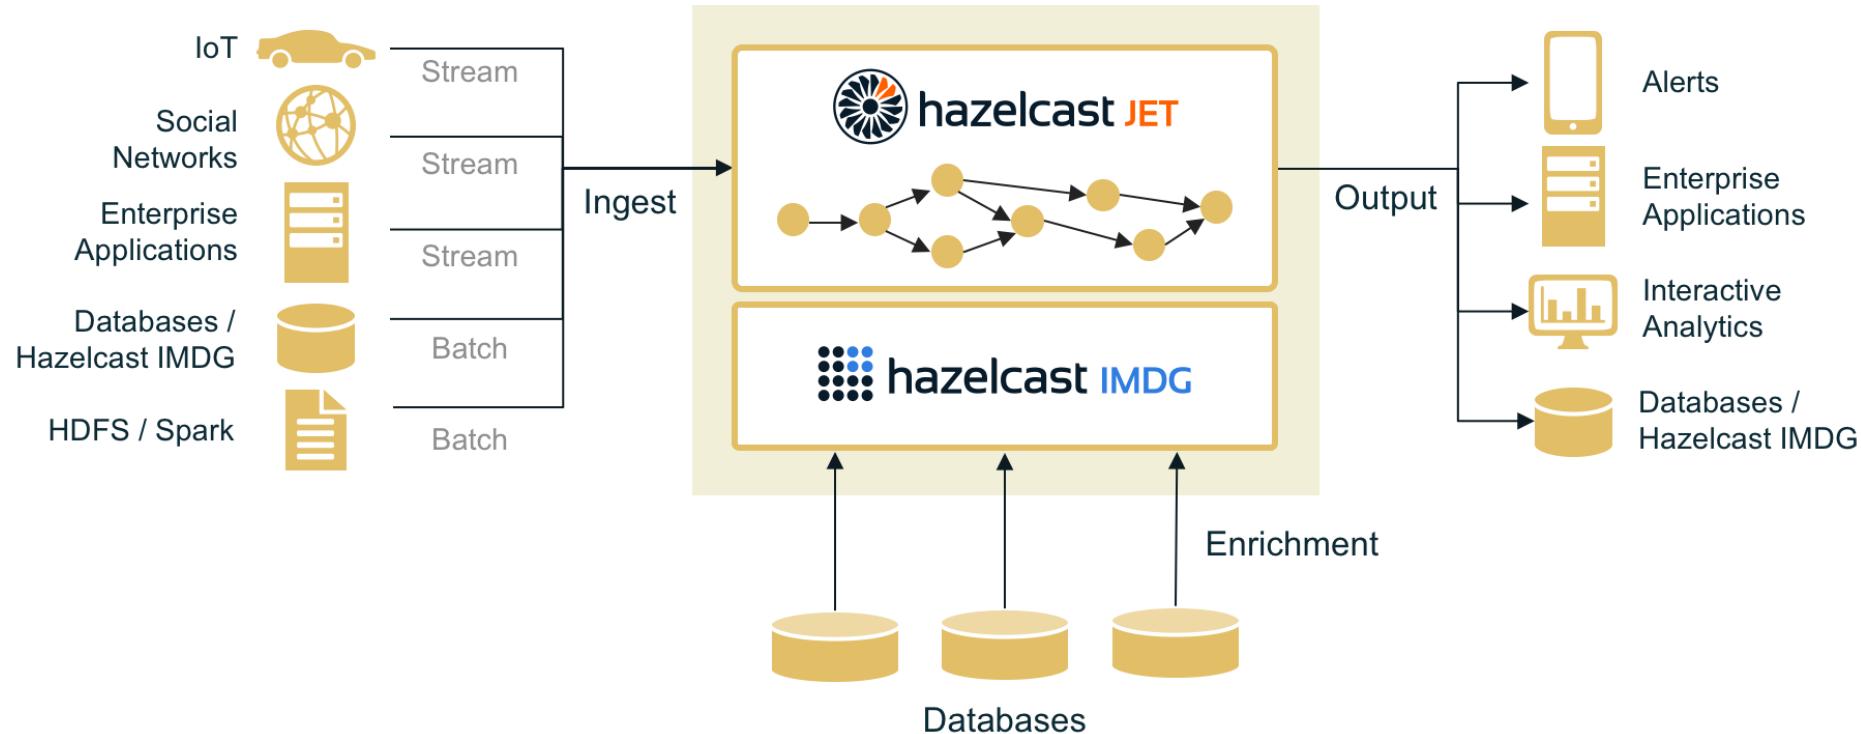
**hazelcast®**



@nicolas\_frankel



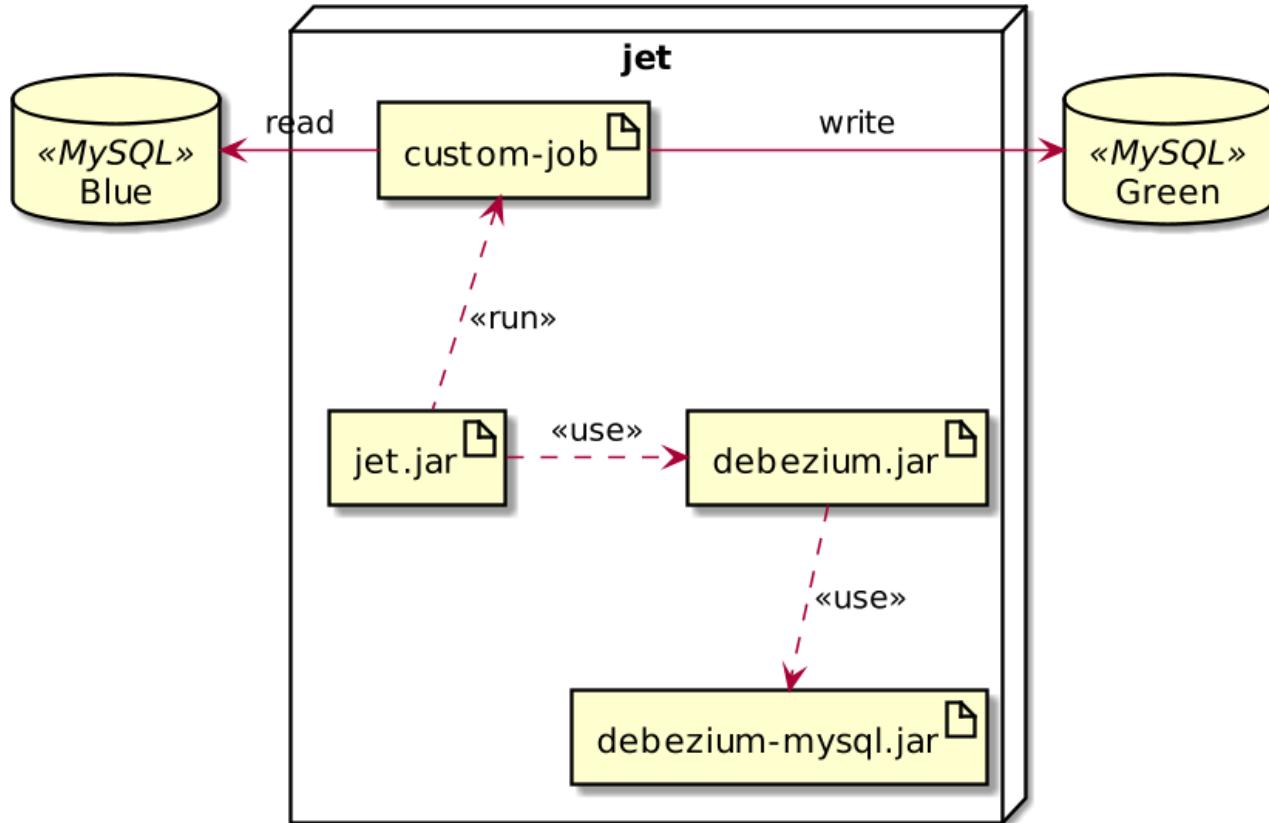
# Hazelcast Jet



@nicolas\_frankel



# Hazelcast Jet & Debezium



@nicolas\_frankel





@nicolas\_frankel



# Takeaways

1. Zero-downtime is within your reach
2. Session replication
3. Change-Data-Capture + Data Streaming for the database



@nicolas\_frankel

# Thanks for your attention!

- <https://blog.frankel.ch/>
- [@nicolas\\_frankel](https://twitter.com/nicolas_frankel)
- <https://bit.ly/zero-downtime>
- <https://slack.hazelcast.com/>
- <https://training.hazelcast.com/>



@nicolas\_frankel